

Measuring Starlink Latency from RIPE Atlas

CHRIS GU, CHRIS YUAN, ETHAN STANLEY, DEVAN FLORES, and CHARLIE MUELLER

The Starlink Network is a collection of satellites in orbit around Earth. The satellites orbit much lower than different artificial satellites and communicate with user terminals on the ground. Our job was to use a set of RIPE Atlas probes connected to the Starlink Network to run measurements on the network. More specifically, we needed to measure latency from user terminals to a destination on the internet across the world.

ACM Reference Format:

Chris Gu, Chris Yuan, Ethan Stanley, Devan Flores, and Charlie Mueller. 2023. Measuring Starlink Latency from RIPE Atlas. 1, 1 (May 2023), 8 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

As mentioned above, we took measurements over the Starlink satellite network. These satellites stay close to Earth (around a few hundreds of kilometers). Since they stay so close to Earth the satellites are believed to give high throughput and low latency between different terminals. In this way different user terminals can send data to each other via the satellites. In this project we wanted to measure the latency from user terminals to a site on the Internet. To do this, it is important to know that Starlink also consists of ground stations on Earth, which serve as a connection to the Internet's fiber infrastructure. The satellites can send signals between user terminals and different ground stations. In this way, we can measure the latency from different Starlink user terminals to the Internet (user terminals send signal to satellites which bounce the signal to the nearest ground station for a site on the Internet).

How did we take these measurements? The company RIPE Atlas has a set of probes that can run and collect measurements on a given network. These probes are given to volunteers who connect a probe to their network. So, we considered the set of RIPE probes connected to the Starlink network through a Starlink user terminal. Furthermore, each probe has an ID number and some other helpful information such as the country it is located in. The measurement we are interested in is ping. Ping is the time for a user to send packets to a destination, and then the time needed to receive a response from the destination. Using ping we can get a sense of the latency of a network since ping measures RTT (round trip time). Our goal for this project was to measure the RTT time from Starlink user terminals to google.com over the course of a day using a set of RIPE probes.

2 CREDIT CALCULATION

It costs credits on the RIPE Atlas site to use probes to collect measurement data. Cost for ping measurements is calculated based on the total number of measurements and the amount and size of packets per ping.

$$\text{Unit Cost} = N * (\text{int}(S/1500) + 1) \quad (1)$$

Authors' address: Chris Gu; Chris Yuan; Ethan Stanley; Devan Flores; Charlie Mueller.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2023 Association for Computing Machinery.

XXXX-XXXX/2023/5-ART \$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

Where N is the number of packets and S is the size in octets. After doing some reading on the RIPE Atlas site we decided to send two packets for each ping measurement, giving us two RTT's for each measurement. We believed that by averaging these results we would get less variable data since after testing some of the pings we saw high variance in the RTT for different packets of the same ping. The project required us to send a ping from a probe every 15 minutes for the entirety of 24 hours. This gives us 96 pings per probe, and 192 packets per probe over the day. As will be explained below, we ended up finding 56 RIPE probes connected to the Starlink network. $192 * 56$ gave us a final total of 10,752 packets, or 10,752 credits that we needed from the course staff to get the data.

3 FINDING PROBES

As explained above, to take these measurements we needed to find all the RIPE Atlas probes that are connected to Starlink user terminals. Each probe has an ID number as specified by the RIPE Atlas API, which also offered a handful of functions helpful in finding probes and creating measurements. One particular API function we used was `GET /api/v2/probes/`, which returns a list of RIPE Atlas probes. In the API call we added in a parameter `asn_v4=14593`, which causes the API to only return probes connected to the system with `ASN = 14593`. We found out that `ASN 14593` corresponds to the Starlink network through an online search. This enabled us to use that API call to get a list of RIPE Atlas probes connected to the Starlink network. We then also used a second parameter in the API call, `status=1`, which caused the API to only return probes that are active. We then called the API, and turned the response into JSON. This allowed us to iterate through the returned list of probes and create a new list containing only the probe ID and country codes for each probe. This gave us what we needed to begin measuring: a list of all probe IDs corresponding to RIPE Atlas probes that are active and connected to the Starlink network. After running everything, we had a list of 56 RIPE probes. Below is a snippet of the code we wrote to find the needed probes.

```
#PART 1
def fetch_probe_list():
    #Function to return list of active Starlink Probes
    starlink_probe_list = []
    api_endpoint = "https://atlas.ripe.net/api/v2/probes/?asn_v4=14593&status=1"

    params = {
        "asn_v4": 14593, #Starlink ASN number
        "status": 1      #Connected Probes
    }

    response = requests.get(api_endpoint, params=params)
    json_data = response.json()

    for probe in json_data["results"]:
        starlink_probe_list.append({"id": probe["id"], "country_code": probe["country_code"]})

    print("Number of Probes: ", len(starlink_probe_list))
    return starlink_probe_list
```

Fig. 1. Code used to get a list of RIPE Atlas Probes connected to Starlink.

4 CREATING MEASUREMENTS AND GETTING DATA

To create a measurement, we can send a post request to <https://atlas.ripe.net/api/v2/measurements> with the specifications of our measurement. Using a definitions object in our header, we can specify what type of measurement we want (ping in this case), the time, interval, packet count and size, and etc. We can also define with probes to measure on explicitly in the body. Using the probes retrieved from [GET /api/v2/probes/](https://atlas.ripe.net/api/v2/probes/), we set the measurement to collect from those 57 probes by listing the probe IDs. Finally, we can send this POST request with our an API key generated by our RIPE Atlas account.

```
definitions = {  
    "definitions": [  
        {  
            "type": "ping",  
            "target": "bing.com",  
            "af": 4,  
            "description": "test trial",  
            "start_time": cur_time + 300,  
            "end_time": cur_time + 86700,  
            "interval": 900,  
            "packets": 2,  
            "size": 64,  
        }  
    ],  
    "probes" : [  
        {  
            "requested": 57,  
            "type": "probes",  
            "value": "35681, 19983, 24742, 52918, 1004453, 1004828, 1001356, 35042, 55492, 60510, 26834,"  
        }  
    ]  
}  
  
def create_measurement(defs):  
    api_endpoint = "https://atlas.ripe.net/api/v2/measurements"  
    params={"key": "XXXXXXXXXXXXXXXXXXXX"}  
  
    x = requests.post(api_endpoint, json = defs, params=params)  
    json_data = x.json()  
    return json_data
```

Fig. 2. Code used to create a ping measurement

Retrieving results just requires the measurement ID which is returned as a result after posting a new measurement.

```
def get_results(id):
    api_endpoint = "https://atlas.ripe.net/api/v2/measurements/" + id + "/results"
    x = requests.get(api_endpoint)
    json_data = x.json()
    return json_data
```

Fig. 3. Code used to retrieve a measurement.

5 PROCESSING DATA AND PLOTTING

We processed the data to remove outliers using standard deviation of response times based on the probe. We plotted the data based on location of the probe as well as probe ID. Because Starlink heavily depends on ground station, plotting based on location will give more insightful results. Probes that are in the same approximate area also experience similar weather conditions and make connections to a close if not the same satellite. These factors can affect the internet connection and therefore should be account for in plotting.

```
import matplotlib.pyplot as plt

# Create a plot with multiple lines
fig, ax = plt.subplots(figsize=(15,6))
for probe, values in probe_data.items():
    timestamps = [item['timestamp'] - 810897 for item in values if item['timestamp'] - 810897 < (86400+1.682e9)]
    avgs = [item['avg'] for item in values if item['timestamp'] - 810897 < (86400+1.682e9)]
    ax.plot(timestamps, avgs, label=probe)

# Set labels and title
ax.set_xlabel("Time (s)")
ax.set_ylabel("RTT avg (ms)")
ax.set_title("RTT avg vs Time for Multiple Probes")

# Add a legend
ax.legend()

# Show the plot
plt.show()
```

Fig. 4. Code for plot generation

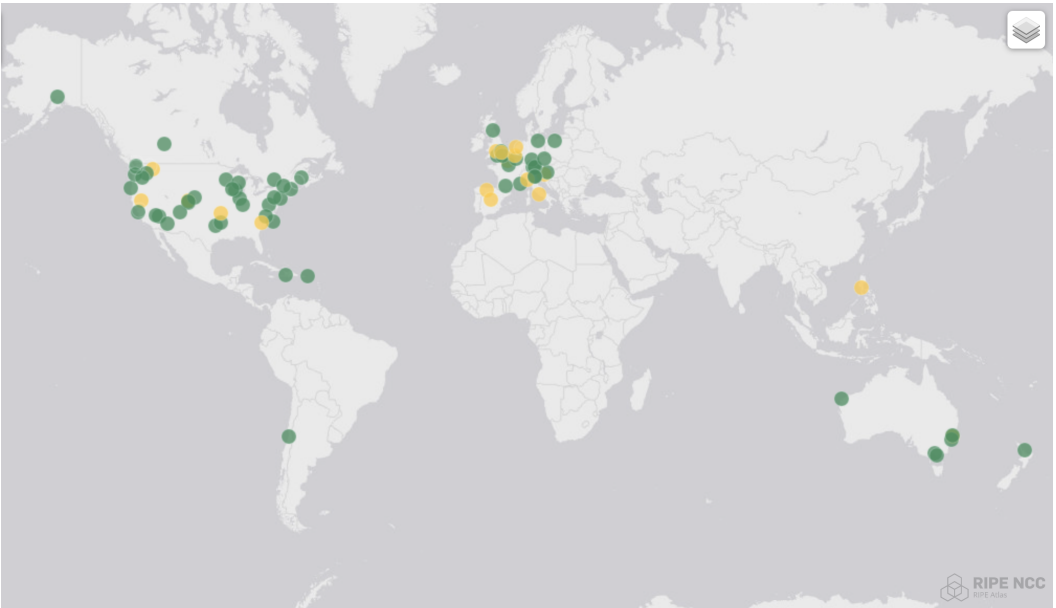


Fig. 5. Starlink probes overlayed onto World Map

We used the python library matplotlib to generate plots. Timestamps of our data is subtracted by 810897 to zero the charts. All plots are round-trip time average over UTC time.

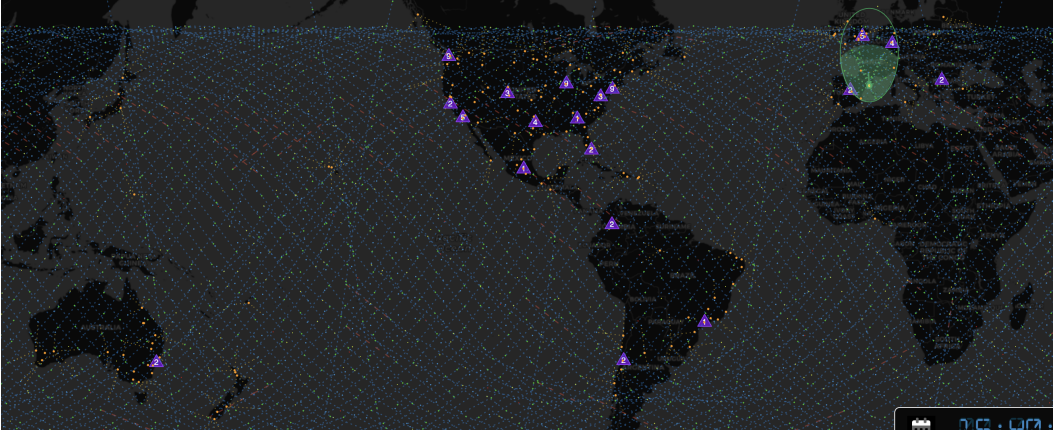


Fig. 6. Purple Dots represent ground stations

6 ANALYZING RESULTS

Figures 6 through 9 display the average round trip time (RTT) ping of various probes over the course of our 24-hour collection period. Figure 6 includes data collected from all probes, which demonstrates the high degree of variance across the different probe locations globally. However, we can refine our analysis with more specific plots and make more specific conclusions regarding how geographic location or time of day may affect RTT.

For example, we compare in Figure 7 the RTT of probe 61113 in Alaska and Probe 26834 in South America, near Santiago, Chile. We see that there is a clear performance difference, with the Alaskan probe experiencing large ping with high variance across the board. In contrast, the Chilean probe has consistently low RTT, around 45ms in the 24h cycle. These are two probes both outside of regions where most RIPE Atlas probes are located, but from Figure 5, we see that there is a ground station located near the Chilean probe while the Alaskan probe exists far from the nearest station. This likely plays a major role in the disparity plotted; regardless of satellite constellation placement, being far from a ground station would require a consistently longer network path through more satellites.

Figure 8 further demonstrates the dependence of RTT ping on the probe distance to a ground station. This plot contains the RTT averages for all RIPE Atlas probes located in Australia. Probe 52918 has consistently the worst performance across the 24-hour period, because it is located in Exmouth, a town on the western coast, while the only ground station in Australia is located on the Southeast corner of country.

In Figure 9, we plot the RTT averages of American RIPE Atlas probes, labelling the ones located in rural areas in cyan and the urban ones in red. We do not see a clear disparity between the two groups, although there are a couple rural outlier probes with highly variant RTT, as seen by the cyan spikes. However, Figure 5 shows that there Starlink ground stations pretty equally dispersed across the continental United States, so it would make sense that both rural and urban areas see mostly similar RTT performances.

By and large, we see that probe distance to Starlink ground stations is the most important factor when it comes to influencing Round Trip Time ping.




Fig. 7. RTT Avg over Time for all probes

7 POTENTIAL IMPLICATIONS

Ultimately, our research implies that Starlink is meant as an alternative to standard broadband connections, instead of a replacement. In most scenarios, Starlink performs optimally and gives users a stable and reliable connection. However, there are several outliers that show significant issues. We believe this is due to the very nature of satellite connections and the various interruptions that can happen when transmitting over satellite. In practice, many Starlink customers complain about the unreliable nature of the technology. Everyday objects such as "trees, birds, and buildings" can cause major interruptions in signal [1]. Starlink themselves even say that the technology was designed for areas "where connectivity has been unreliable or completely unavailable" [2]. Although Starlink is a promising technology, we are very far away from completely replacing broadband.

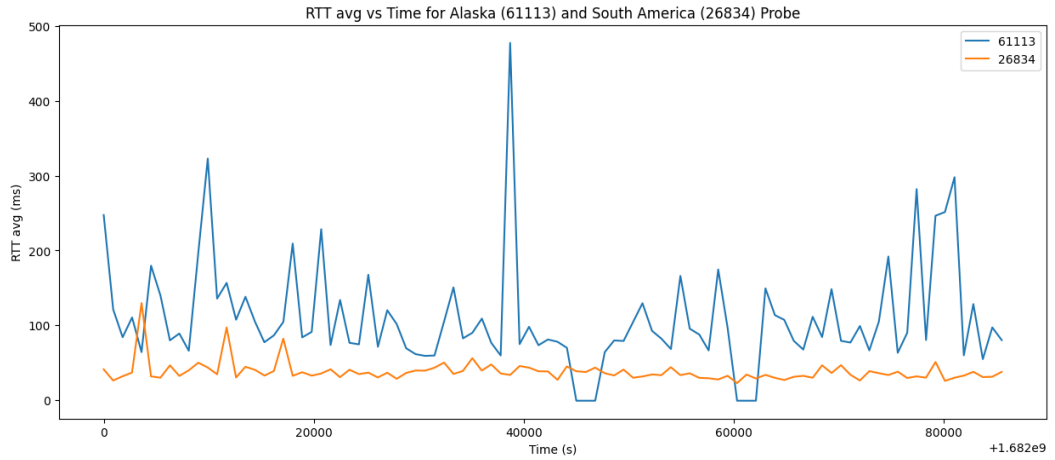


Fig. 8. RTT Avg over Time for Alaska (61113) and South America (26834) probes

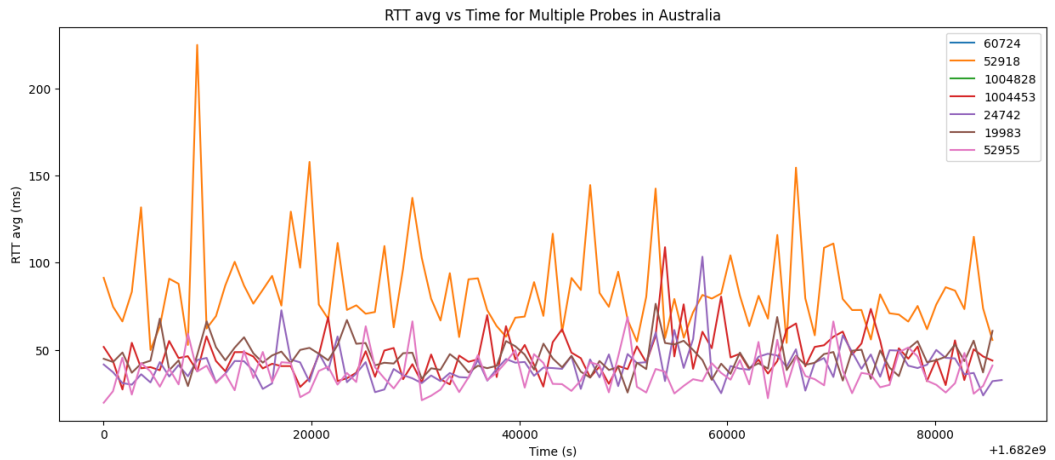


Fig. 9. RTT Avg over Time for Australian probes

REFERENCES

- [1] Nilay Patel. 2021. Starlink review: broadband dreams fall to Earth. (14 May 2021). Retrieved May 9, 2023 from <https://www.theverge.com/22435030/starlink-satellite-internet-spacex-review>
- [2] Starlink. 2023. *Roam*. Retrieved May 9, 2023 from <https://www.starlink.com/roam>

A GITHUB REPOSITORY

Please access our github repository with relevant code here: <https://github.com/cgu2020/project-starlink>

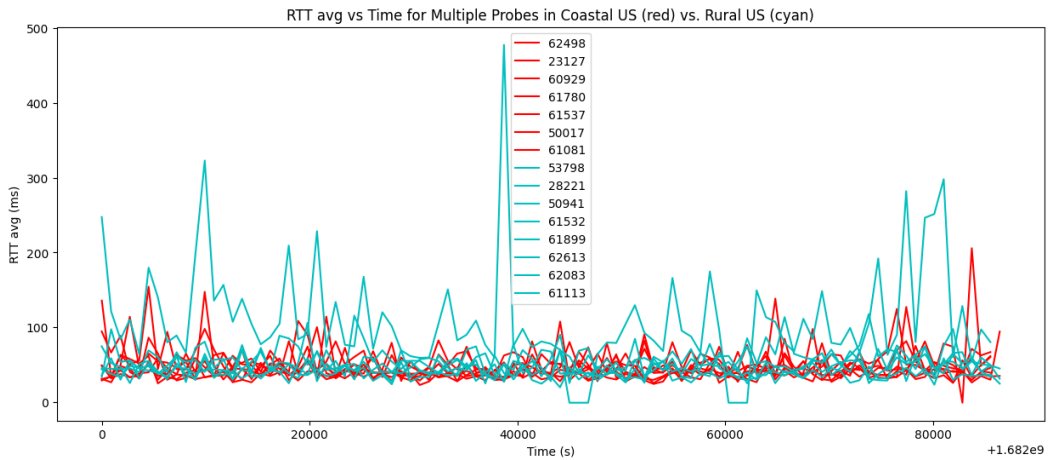


Fig. 10. RTT Avg over Time for Coastal vs Rural America probes